

1. High-Level Ask

1.1 Objective

Build a **GPU-accelerated AprilTag detection engine** on Jetson Orin NX that can achieve **effective ~120 FPS** using 720p camera input for FRC-style localization, with performance substantially exceeding the current CPU-only implementation.

1.2 Scope

- Single camera initially (IMX477 / AR0234 @ 1280×720), scalable to multiple cameras later.
- Detection of **FRC AprilTags** (e.g. TAG_36H11) with:
 - Reliable detection up to typical FRC distances (3–8 m+),
 - Latency low enough to support fast swerve control and pose correction.
- GPU handles:
 - Raw → Gray conversion
 - Undistortion + decimation
 - Gradient / edge / quad candidate generation
 - Tag decode
 - PnP (3D pose)
- CPU handles:
 - Camera capture to pinned host memory
 - High-level scheduling / control
 - Consuming final poses (e.g. pushing to NetworkTables / Pose Estimator).

1.3 Success Criteria

- **Performance:**
 - Baseline: >40 FPS full-frame AprilTag detection at 720p input with decimated working resolution (e.g. 640×360).
 - Target: effective **120 Hz pose updates** using full-frame + ROI strategy.
- **Quality:**
 - Detection rate and pose accuracy comparable to apriltag3 CPU implementation at equivalent resolution.
- **Scalability:**
 - Architecture supports extension to up to 4 cameras with modest rework.
- **Robustness:**
 - Handles camera disconnects, GPU errors, and recovers gracefully.

1.4 Constraints

- Platform: **Jetson Orin NX**, using CUDA.
 - Language: C++ (core) + Python wrapper (for dev/testing).
 - Time budget: this is a non-trivial R&D project; we'll stage it in phases.
-

2. System Architecture Overview

2.1 Top-Level Data Flow

Per frame:

1. **Camera Capture (CPU)**
IMX477 / AR0234 → raw frame into pinned host buffer (BAYER_RG8 / GRAY8).
2. **Upload & Preprocess (GPU)**

- Async upload to GPU.
- Fused kernel: raw → Gray8 → undistort → decimate → working image.

3. AprilTag Detection (GPU)

- Kernel 1: gradients + edge maps.
- Kernel 2: quad candidate extraction.
- Kernel 3: decode + PnP.

4. Results Return (CPU)

- Copy small detection list (IDs + poses) back to CPU.
- Feed into robot code / NT.

2.2 Components

1. Camera HAL ([ICameraDevice](#))
 2. Frame Manager (CPU) – double/triple buffering of pinned host frames.
 3. GPU Context ([GpuContext](#)) – streams + device buffers.
 4. Image Preprocessor ([ImagePreprocessor](#)) – fused kernel(s) for raw→Gray.
 5. AprilTag GPU Detector ([AprilTagCudaDetector](#)) – detection kernels.
 6. Scheduler/Pipeline ([CameraPipeline](#) / [MultiCameraSystem](#)) – orchestrates frames.
 7. Persistent Kernel Framework (optional later) – eliminates per-frame launch overhead.
 8. Python Wrapper ([c5b_apriltag](#)) – for experimentation.
-

3. Component Design & Implementation Plan

3.1 Camera HAL & Frame Manager

Goal: Fast, predictable camera→CPU→GPU path with minimal overhead.

3.1.1 ICameraDevice

Already conceptually defined:

```
class ICameraDevice {
public:
    virtual ~ICameraDevice() = default;
    virtual const CameraStaticInfo& staticInfo() const = 0;
    virtual const CameraRuntimeConfig& runtimeConfig() const = 0;
    virtual const CameraCalibration& calibration() const = 0;

    virtual bool grabRaw(uint8_t* dst,
                         int dstStrideBytes,
                         RawFrameFormat& outFormat) = 0;
};
```

Implementation: `UvcCameraDevice` using V4L2/OpenCV.

- Configure for **1280×720** at desired FPS.
- Prefer `BAYER_RG8` or `GRAY8` for AR0234, `BAYER_RG8` or `BGR8` for IMX477.

3.1.2 Frame Manager (CPU)

Allocate **N pinned buffers** (2–3) using `cudaHostAlloc`:

```
struct HostFrameBuffer {
    uint8_t* ptr;
    int      stride;
};

class FrameManager {
public:
    FrameManager(int width, int height, PixelFormat fmt, int numBuffers);
    ~FrameManager();
```

```
    HostFrameBuffer acquireForCapture(); // returns free buffer
    void releaseAfterGPU(HostFrameBuffer&); // after upload
};
```

Camera thread:

- `acquireForCapture()` → `grabRaw()` → enqueue for GPU.
-

3.2 GPU Context & Buffers

Goal: Preallocate all GPU memory and streams for predictable high FPS.

3.2.1 GpuContext extensions

```
class GpuContext {
public:
    explicit GpuContext(int deviceIndex = 0);
    ~GpuContext();

    cudaStream_t createStream(const std::string& name);
    void         destroyStream(cudaStream_t stream);

    RawGpuImage allocRawImage(int width, int height, PixelFormat fmt,
    int numBuffers);
    GpuImage     allocGrayImage(int width, int height, int numBuffers);

    void uploadRawAsync(const HostFrameBuffer& host,
                        RawGpuImage& rawGpu,
                        int bufferIndex,
                        cudaStream_t stream);
};
```

- Allocate **rawGpu buffers** and **grayGpu buffers** in arrays to support double/triple buffering.
-

3.3 Image Preprocessor – Fused Kernel

Goal: One pass: raw → gray → undistort → decimate (→ optionally gradient precompute later).

3.3.1 API

```
struct PreprocessParams {
    bool undistort;
    bool decimate;
    float decimateFactor; // e.g. 2.0f
};

class ImagePreprocessor {
public:
    ImagePreprocessor(GpuContext& ctx,
                      const CameraCalibration& calib,
                      int inputWidth,
                      int inputHeight,
                      float decimateFactor);

    // Fused preprocess: rawGpu -> grayGpu
    void preprocessAsync(const RawGpuImage& rawGpu,
                         GpuImage& grayGpu,
                         const PreprocessParams& params,
                         cudaStream_t stream);

private:
    // Optional: precomputed undistortion map, stored on device
    GpuImage undistortMap_; // or a custom struct for float2 mapping
};
```

3.3.2 Implementation steps

1. Startup:

- Precompute undistort map for 1280×720 (or working resolution) if `undistort = true`.
- Store mapping from output pixel to input pixel (`float2`) in a device buffer.

2. Kernel: `rawToGrayUndistortDecimateKernel`

- Grid: 2D, covering **output (decimated) resolution**, e.g. 640×360.
- Per-thread (for each output pixel):
 - Compute or lookup undistorted input coordinates (`x_in`, `y_in`).
 - For `BAYER_RG8`:
 - Sample Bayer pattern, compute intensity (e.g. average of nearby green pixels).
 - For `GRAY8`:
 - Direct sample.
 - Write Gray8 value to `grayGpu`.
- Optimize for:
 - Coalesced reads from raw image.
 - Coalesced writes to gray image.
 - If using map, compute once on CPU/GPU and reuse.

Performance target: << 1 ms per frame at 640×360.

3.4 AprilTag GPU Detector – Kernels

We implement `AprilTagCudaDetector` as a sequence of optimized kernels over the decimated Gray image.

3.4.1 API

```
class AprilTagCudaDetector {  
public:  
    AprilTagCudaDetector(const AprilTagDetectorConfig& cfg,  
                         const CameraCalibration& calib,  
                         GpuContext& gpuCtx,
```

```

        int workWidth,
        int workHeight);

    std::vector<TagDetection3D> detect(const GpuImage& grayImg,
                                         cudaStream_t stream,
                                         const float cameraToRobot[16] =
                                         nullptr);
}

```

Internally:

- Uses:
 - Device buffers for gradient/edge maps.
 - Device buffer for quad candidates.
 - Device buffer for final detections (TagDetection3D).

3.4.2 Kernel 1 – Gradient & Edge Map

Input: Gray8 image (decimated).

Output:

- Gradient magnitude (or binary edge map).
- Possibly orientation.

Kernel design:

- Block: 16×16 or 32×8 threads.
- Use shared memory tiles with halo for Sobel filter.
- Steps per pixel:
 - Load neighbor pixels into shared memory.
 - Compute Gx, Gy; magnitude = |Gx|+|Gy| or sqrt.

- Threshold to produce edge map.
-

3.4.3 Kernel 2 – Quad Candidate Extraction

Goal: Identify quadrilateral edge clusters as candidate tags.

Strategy:

- Divide image into tiles (e.g., 32×32).
- Each block:
 - Uses shared memory to analyze edge pixels in tile.
 - Builds line segments or short chains of edge pixels.
 - Groups segments into 4-sided shapes with geometric heuristics:
 - roughly rectangular,
 - roughly convex,
 - size above a minimum pixel area.
- Accepted candidates:
 - Append to a global `QuadCandidate` array using atomic add to maintain count.
 - Cap count to some max (e.g., 256 or 512) to bound worst-case.

Data structure:

```
struct QuadCandidate {  
    float corners[4][2]; // x,y pixel positions  
    float score;         // heuristic quality  
};
```

3.4.4 Kernel 3 – Decode + PnP

One block per quad candidate:

- Normalize quad to tag coordinate frame (homography).
- Sample interior bit grid (e.g., 6×6 / 8×8), compute ID + Hamming.
- Reject if Hamming too high or margin too low.
- For valid tags:
 - Compute PnP (pose) on GPU:
 - Use known tag size + four corners.
 - Use simplified PnP (e.g., P3P + refinement or EPnP).
 - If `cameraToRobot` present:
 - Apply 4×4 transform to convert to robot frame.

Write results into a `TagDetection3D` array on device, with atomic increment of detection count.

Finally:

- Copy results to host (`cudaMemcpyAsync`) – very small data.
-

3.5 Persistent Kernel (Phase 2 / Advanced)

To approach 120 FPS, we want to **reduce per-frame launch overhead**.

3.5.1 Concept

- Launch a **single persistent kernel** with a small number of SMs reserved.
- Kernel loops over a work queue of frames:
 - For each frame: run preprocess + detect pipeline.
- CPU only:

- Writes frame info (buffer indices) to queue in mapped memory.
- Signals kernel via atomic flag.
- Polls or waits for result.

3.5.2 Work Queue Structure

```
struct FrameJob {
    int rawIndex;
    int grayIndex;
    int jobId;
    int status; // 0=pending, 1=processing, 2=done
};

struct FrameQueue {
    FrameJob jobs[MAX_JOBS];
    int head;
    int tail;
};
```

Persistent kernel:

- Dequeues a job, runs full pipeline (preprocess + detection), marks job as done.

CPU:

- Enqueues a job after capture/upload, waits for `status==2`, then copies detections.

This is an advanced step; you can start with normal per-frame kernels, then replace them with a persistent kernel once everything works.

3.6 Multi-Rate / ROI Strategy (Practical Route to 120 Hz)

True full-frame heavy search at exactly 120 FPS may still be overkill. A practical architecture:

1. **Full-frame detection at 30–40 FPS**, using the full pipeline described above.

2. ROI-based detection at 120 FPS:

- Maintain expected tag locations from last frame + robot motion.
- For intermediate frames:
 - Only run detection kernels in small ROIs around these predicted tag positions.
- That massively reduces per-frame workload.

Architecturally:

- `AprilTagCudaDetector` supports:
 - A “full-frame” mode (grid spans entire image).
 - An “ROI list” mode (kernel only processes bounding boxes).
-

4. CameraPipeline & Scheduler

4.1 CameraPipeline (Single Camera)

```
class CameraPipeline {  
public:  
    CameraPipeline(const CameraPipelineConfig& cfg, GpuContext&  
gpuCtx);  
    ~CameraPipeline();  
  
    std::vector<TagDetection3D> processOnce(); //  
non-persistent-kernel mode  
  
private:  
    GpuContext& gpuCtx_;  
    std::unique_ptr<ICameraDevice> camera_;  
    FrameManager frameMgr_;  
    ImagePreprocessor preproc_;  
    AprilTagCudaDetector detector_;
```

```

    cudaStream_t          stream_ ;

    int currentBufferIndex_ ; // for double/triple buffering
};


```

Flow in `processOnce()`:

1. `HostFrameBuffer buf = frameMgr_.acquireForCapture();`
2. `camera_->grabRaw(buf.ptr, buf.stride, rawFormat);`
3. `gpuCtx_.uploadRawAsync(buf, rawGpu, idx, stream_);`
4. `preproc_.preprocessAsync(rawGpu, grayGpu, params, stream_);`
5. `detections = detector_.detect(grayGpu, stream_, cameraToRobot);`
6. `frameMgr_.releaseAfterGPU(buf);`
7. Return detections.

Later, for persistent kernel mode, `processOnce()` will instead:

- Post a job into the persistent kernel queue and wait for completion.
-

5. Python Wrapper for Dev

Expose a minimal high-level API:

```

import c5b_apriltag as c5b

gpu = c5b.GpuContext()
cfg = c5b.load_config("camera.yaml")
cam = c5b.Camera(gpu, cfg)

while True:

```

```
dets = cam.process_once()  
for d in dets:  
    print(d.id, d.pose)
```

Later add:

- Multi-camera support
- ROI mode
- Debug: visualize candidate edges/quads.